

PSET0 + Python & Linear Algebra Review

CS 231A
04/04/2025

A bit about me!

- Coterm CS + Upcoming CS PhD!
- Work in SVL with Professors Fei-Fei Li and Ehsan Adeli
- Research interests in 3D computer vision for medical applications
- How can we extract insights about a 3D world from monocular video?
 - Monocular depth estimation
 - Scene reconstruction
 - Scene understanding
- Office Hours: Thursdays 2-3 PM in CoDa B43 (check Canvas for Zoom)



Outline

- Python Introduction
- Linear Algebra and NumPy
- PSET0

Outline

- **Python Introduction**
- Linear Algebra and NumPy
- PSET0

Python



High-level, interpreted programming language.

Python will be used in all the homeworks and recommended for the project.

We'll cover some basics today.



Why Python?

- Python is high-level.

JAVA

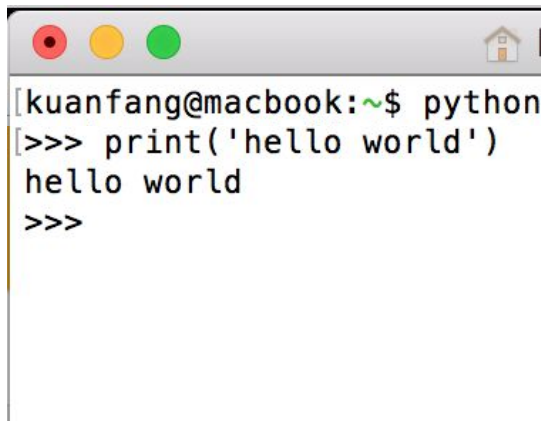
```
public class Main {  
    public static void main(String[] args) {  
        System.out.println("hello world");  
    }  
}
```

PYTHON


```
print('hello world')
```

Why Python?

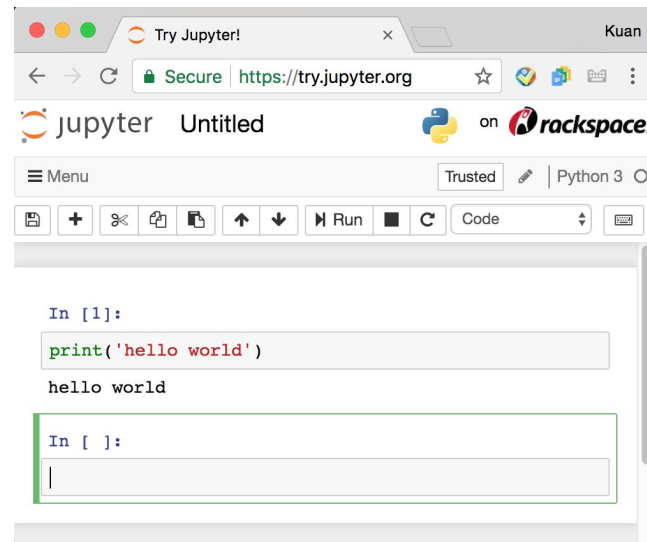
- Python is accessible.



Interpreter/Terminal



IDE





Jupyter Notebook



Why Python?

- Python has many many awesome packages.



How to Set up Python?

1. Follow this guide: <https://wiki.python.org/moin/BeginnersGuide/Download>
2. Choose your favourite editor or IDE:
 - a. Sublime
 - b. Vim
 - c. VSCode
 - d. Spyder
 - e. PyCharm
 - f. Jupyter Notebook
 - g. Colab
 - h. ...

Basic Python Review


Google Colab [Here](#)

Outline

- Python Introduction
- **Linear Algebra and NumPy**
- PSET0


Why use Linear Algebra in Computer Vision?

- As you've seen in lecture, using linear algebra is necessary to represent many quantities, e.g. 3D points on a scene, 2D points on an image.
- Transformations of 3D points with 2D points can be represented as matrices.
- Images are literally matrices filled with numbers (as you will see in PSET0).



Vector Review

$$\mathbf{v} = (x_1, x_2)$$



Magnitude: $\|\mathbf{v}\| = \sqrt{x_1^2 + x_2^2}$

If $\|\mathbf{v}\| = 1$, \mathbf{v} is a UNIT vector

$$\frac{\mathbf{v}}{\|\mathbf{v}\|} = \left(\frac{x_1}{\|\mathbf{v}\|}, \frac{x_2}{\|\mathbf{v}\|} \right) \text{ is a unit vector}$$

Orientation: $\theta = \tan^{-1}\left(\frac{x_2}{x_1}\right)$

Vector Review

$$\mathbf{v} + \mathbf{w} = (x_1, x_2) + (y_1, y_2) = (x_1 + y_1, x_2 + y_2)$$

$$\mathbf{v} - \mathbf{w} = (x_1, x_2) - (y_1, y_2) = (x_1 - y_1, x_2 - y_2)$$

$$a\mathbf{v} = a(x_1, x_2) = (ax_1, ax_2)$$

Matrix Review

$$A_{n \times m} = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1m} \\ a_{21} & a_{22} & \dots & a_{2m} \\ \vdots & \vdots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nm} \end{bmatrix}$$



Pixel's intensity value

$$\text{Sum: } C_{n \times m} = A_{n \times m} + B_{n \times m} \quad c_{ij} = a_{ij} + b_{ij}$$

A and B must have the same dimensions!

$$\text{Example: } \begin{bmatrix} 2 & 5 \\ 3 & 1 \end{bmatrix} + \begin{bmatrix} 6 & 2 \\ 1 & 5 \end{bmatrix} = \begin{bmatrix} 8 & 7 \\ 4 & 6 \end{bmatrix}$$

Matrices and Vectors in Python (NumPy)



```
import numpy as np
```

An optimized, well-maintained scientific computing package for Python.

As time goes on, you'll learn to appreciate NumPy more and more.

np.ndarray: Matrices and Vectors in Python

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
import numpy as np

M = np.array([[1, 2, 3],
              [4, 5, 6],
              [7, 8, 9]])

v = np.array([[1],
              [2],
              [3]])
```

np.ndarray: Matrices and Vectors in Python

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
print(M.shape)    # (3, 3)
print(v.shape)    # (3, 1)
```

np.ndarray: Matrices and Vectors in Python

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
print(v + v)
```

```
[[2]  
[4]  
[6]]
```

```
print(3 * v)
```

```
[[3]  
[6]  
[9]]
```

Other Ways to Create Matrices and Vectors

NumPy provides many convenience functions for creating matrices/vectors.

```
a = np.zeros((2,2)) # Create an array of all zeros
print a             # Prints "[[ 0.  0.]
                    #           [ 0.  0.]]"

b = np.ones((1,2))  # Create an array of all ones
print b             # Prints "[[ 1.  1.]]"

c = np.full((2,2), 7) # Create a constant array
print c             # Prints "[[ 7.  7.]
                    #           [ 7.  7.]]"

d = np.eye(2)        # Create a 2x2 identity matrix
print d             # Prints "[[ 1.  0.]
                    #           [ 0.  1.]]"

e = np.random.random((2,2)) # Create an array filled with random values
print e              # Might print "[[ 0.91940167  0.08143941]
                    #           [ 0.68744134  0.87236687]]"
```

Matrix Indexing

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$


```
print(M)
```

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

```
print(M[:2, 1:3])
```

```
[[2 3]
 [5 6]]
```

Dot Product




$$v \cdot w = (x_1, x_2) \cdot (y_1, y_2) = x_1 y_1 + x_2 y_2$$

The inner product is a **SCALAR!**

$$v \cdot w = (x_1, x_2) \cdot (y_1, y_2) = \|v\| \cdot \|w\| \cos \alpha$$

$$\text{if } v \perp w, \quad v \cdot w = ? = 0$$


Cross Product



$$u = v \times w$$

The cross product is a **VECTOR!**

Magnitude: $\|u\| = \|v \times w\| = \|v\| \|w\| \sin \alpha$



Orientation:

$$u \perp v \Rightarrow u \cdot v = (v \times w) \cdot v = 0$$

$$u \perp w \Rightarrow u \cdot w = (v \times w) \cdot w = 0$$

$$\text{if } v \parallel w \rightarrow u = 0$$

Cross Product

$$\mathbf{i} = (1,0,0) \quad \|\mathbf{i}\| = 1 \quad \mathbf{i} = \mathbf{j} \times \mathbf{k}$$

$$\mathbf{j} = (0,1,0) \quad \|\mathbf{j}\| = 1 \quad \mathbf{j} = \mathbf{k} \times \mathbf{i}$$

$$\mathbf{k} = (0,0,1) \quad \|\mathbf{k}\| = 1 \quad \mathbf{k} = \mathbf{i} \times \mathbf{j}$$

$$\begin{aligned} \mathbf{u} &= \mathbf{v} \times \mathbf{w} = (x_1, x_2, x_3) \times (y_1, y_2, y_3) \\ &= (x_2 y_3 - x_3 y_2) \mathbf{i} + (x_3 y_1 - x_1 y_3) \mathbf{j} + (x_1 y_2 - x_2 y_1) \mathbf{k} \end{aligned}$$

Matrix Multiplication

$$A_{n \times m} = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1m} \\ a_{21} & a_{22} & \dots & a_{2m} \\ \vdots & \vdots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nm} \end{bmatrix} \mathbf{a}_i$$

$$B_{m \times p} = \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1p} \\ b_{21} & b_{22} & \dots & b_{2p} \\ \vdots & \vdots & \vdots & \vdots \\ b_{m1} & b_{m2} & \dots & b_{mp} \end{bmatrix} \mathbf{b}_j$$

Product:

$$C_{n \times p} = A_{n \times m} B_{m \times p}$$

$$c_{ij} = \mathbf{a}_i \cdot \mathbf{b}_j = \sum_{k=1}^m a_{ik} b_{kj}$$

A and B must have compatible dimensions!

$$A_{n \times n} B_{n \times n} \neq B_{n \times n} A_{n \times n}$$

Basic Operations - Dot Multiplication

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
print(M.dot(v))
```

```
[[ 14]  
 [ 32]  
 [ 50]]
```

Matrix multiplication in NumPy can be defined as the dot product between a matrix and a matrix/vector.

Basic Operations - Element-wise Multiplication

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
print(np.multiply(M, v))
```

```
[[ 1  2  3]
 [ 8 10 12]
 [21 24 27]]
```

```
print(np.multiply(v, v))
```

```
[[1]
 [4]
 [9]]
```

Broadcasting

- ✓ Between numpy arrays of different shapes
- ✓ Easy and more concise implementations
- ✓ Faster with parallel numerical computations

Broadcasting

<https://numpy.org/doc/1.20/user/theory.broadcasting.html>

```
>>> from numpy import array
>>> a = array([1.0, 2.0, 3.0])
>>> b = array([2.0, 2.0, 2.0])
>>> a * b
array([ 2.,  4.,  6.])
```

```
>>> from numpy import array
>>> a = array([1.0, 2.0, 3.0])
>>> b = 2.0
>>> a * b
array([ 2.,  4.,  6.])
```



Broadcasting




The broadcasting rule

Image	(3d array)	256 x	256 x	3
Scale	(1d array)			3
Result	(3d array)	256 x	256 x	3

A	(4d array)	8 x	1 x	6 x	1
B	(3d array)		7 x	1 x	5
Result	(4d array)	8 x	7 x	6 x	5


Broadcasting

The broadcasting rule




Broadcasting

```
>>> from numpy import array
>>> a = array([[ 0.0,  0.0,  0.0],
...           [10.0, 10.0, 10.0],
...           [20.0, 20.0, 20.0],
...           [30.0, 30.0, 30.0]])
>>> b = array([1.0, 2.0, 3.0])
>>> a + b
array([[ 1.,  2.,  3.],
       [11., 12., 13.],
       [21., 22., 23.],
       [31., 32., 33.]])
```



Broadcasting

```
>>> from numpy import array, newaxis
>>> a = array([0.0, 10.0, 20.0, 30.0])
>>> b = array([1.0, 2.0, 3.0])
>>> a[:,newaxis] + b
array([[ 1.,  2.,  3.],
       [11., 12., 13.],
       [21., 22., 23.],
       [31., 32., 33.]])
```



Broadcasting

Match the dimension of two arrays

Add new axis: `np.newaxis` or `None`

Examples: `arr[None]`, `arr[:, None]`, `arr[:, :, None, :]`,

`arr[..., None]`, `arr[..., None, :]`

Repeat an array: `np.repeat()`, `np.tile()`

Norm

- Informally a measure of the “length” of a vector
- More formally, any measure $f(x)$ that is:
 - Non-negative: $f(x) \geq 0$ for all x .
 - Definite: $f(x) = 0$ iff and only if $x = 0$
 - Homogeneous: $f(tx) = |t|f(x)$
 - Triangle inequality: $f(x + y) \leq f(x) + f(y)$
- There are also norms for matrices like the Frobenius norm:

$$\|A\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}$$


$$\|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$$

$$\|x\|_1 = \sum_{i=1}^n |x_i|$$

$$\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

Orthonormal Basis

= Orthogonal and Normalized Basis



$$\begin{aligned}\mathbf{i} &= (1,0) & \|\mathbf{i}\| &= 1 \\ \mathbf{j} &= (0,1) & \|\mathbf{j}\| &= 1 \\ \mathbf{i} \cdot \mathbf{j} &= 0\end{aligned}$$

$$\mathbf{v} = (x_1, x_2)$$

$$\mathbf{v} = x_1 \mathbf{i} + x_2 \mathbf{j}$$

$$\mathbf{v} \cdot \mathbf{i} = ? = (x_1 \mathbf{i} + x_2 \mathbf{j}) \cdot \mathbf{i} = x_1 1 + x_2 0 = x_1$$

$$\mathbf{v} \cdot \mathbf{j} = (x_1 \mathbf{i} + x_2 \mathbf{j}) \cdot \mathbf{j} = x_1 \cdot 0 + x_2 \cdot 1 = x_2$$

Transpose

Definition:

$$\mathbf{C}_{m \times n} = \mathbf{A}_{n \times m}^T$$

$$c_{ij} = a_{ji}$$

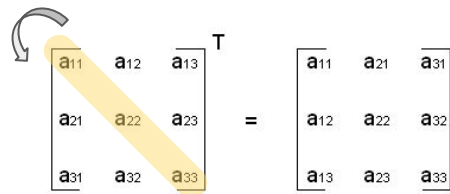
Identities:

$$(\mathbf{A} + \mathbf{B})^T = \mathbf{A}^T + \mathbf{B}^T$$

$$(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$$

If $\mathbf{A} = \mathbf{A}^T$, then \mathbf{A} is *symmetric*

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$


$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}^T = \begin{bmatrix} a_{11} & a_{21} & a_{31} \\ a_{12} & a_{22} & a_{32} \\ a_{13} & a_{23} & a_{33} \end{bmatrix}$$

Basic Operations - Transpose

$$M = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
print(M.T)
```

```
[[1 4 7]
 [2 5 8]
 [3 6 9]]
```

```
print(v.T)
```

```
[[1 2 3]]
```

```
print(M.T.shape)
```

```
print(v.T.shape)
```

```
(3, 3)
```

```
(1, 3)
```

Matrix Determinant

Useful value computed from the elements of a *square* matrix **A**

$$\det [a_{11}] = a_{11}$$

$$\det \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} = a_{11}a_{22} - a_{12}a_{21}$$

$$\det \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} = a_{11}a_{22}a_{33} + a_{12}a_{23}a_{31} + a_{13}a_{21}a_{32} \\ - a_{13}a_{22}a_{31} - a_{23}a_{32}a_{11} - a_{33}a_{12}a_{21}$$

Matrix Inverse

Does not exist for all matrices, necessary (but not sufficient) that the matrix is square

$$\mathbf{A}\mathbf{A}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I}$$

$$\mathbf{A}^{-1} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}^{-1} = \frac{1}{\det \mathbf{A}} \begin{bmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{bmatrix}, \det \mathbf{A} \neq 0$$

If $\det \mathbf{A} = 0$, \mathbf{A} does not have an inverse.

Basic Operations - Determinant and Inverse

$$M = \begin{bmatrix} 3 & 0 & 2 \\ 2 & 0 & -2 \\ 0 & 1 & 1 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
print(np.linalg.inv(M))
```

```
[[ 0.2  0.2  0. ]  
 [-0.2  0.3  1. ]  
 [ 0.2 -0.3 -0. ]]
```

```
print(np.linalg.det(M))
```

```
10.0
```

Matrix Eigenvalues and Eigenvectors

A eigenvalue λ and eigenvector \mathbf{u} satisfies

$$\mathbf{A}\mathbf{u} = \lambda\mathbf{u}$$

where \mathbf{A} is a square matrix.

- Multiplying \mathbf{u} by \mathbf{A} scales \mathbf{u} by λ

Matrix Eigenvalues and Eigenvectors

Rearranging the previous equation gives the system

$$\mathbf{A}\mathbf{u} - \lambda\mathbf{u} = (\mathbf{A} - \lambda\mathbf{I})\mathbf{u} = 0$$

which has a solution if and only if $\det(\mathbf{A} - \lambda\mathbf{I}) = 0$.

- ▶ The eigenvalues are the roots of this determinant which is polynomial in λ .
- ▶ Substitute the resulting eigenvalues back into $\mathbf{A}\mathbf{u} = \lambda\mathbf{u}$ and solve to obtain the corresponding eigenvector.

Basic Operations - Eigenvalues, Eigenvectors

$$M = \begin{bmatrix} 0 & 1 \\ -2 & -3 \end{bmatrix}$$

```
eigvals, eigvecs = np.linalg.eig(M)
```

```
print(eigvals)
```

```
[-1. -2.]
```

```
print(eigvecs)
```

```
[[ 0.70710678 -0.4472136 ]  
 [-0.70710678  0.89442719]]
```

NOTE: Please read the NumPy docs on this function before using it, lots more information about multiplicity of eigenvalues and etc there.

Facts of Eigenvalues

Q1: If an $n \times n$ square matrix is **real**, does it always have n eigenvalues?

Q2: If an $n \times n$ **real** square matrix is **symmetric**, are all its eigenvalues real?

Q3: For any $m \times n$ matrix **A**, does **AA^T** have to be **symmetric**?

Facts of Eigenvalues

Q1: If an $n \times n$ square matrix is **real**, does it always have n eigenvalues?

Yes. But they may not all be real.

Q2: If an $n \times n$ **real** square matrix is **symmetric**, are all its eigenvalues real?

Yes.

Q3: For any $m \times n$ matrix **A**, does **AA^T** have to be **symmetric**?

Yes.

Singular Value Decomposition

$$\begin{matrix} n \\ \boxed{A} \\ m \end{matrix} \approx \begin{matrix} r \\ \boxed{U} \\ m \end{matrix} \times \begin{matrix} r \\ \boxed{\Sigma} \\ r \end{matrix} \times \begin{matrix} n \\ \boxed{V^T} \\ r \end{matrix}$$

Singular values: Non negative square roots of the eigenvalues of $\mathbf{A}^t\mathbf{A}$. Denoted σ_i , $i=1, \dots, n$

SVD: If \mathbf{A} is a real m by n matrix then there exist orthogonal matrices \mathbf{U} ($\in \mathbb{R}^{m \times m}$) and \mathbf{V} ($\in \mathbb{R}^{n \times n}$) such that

$$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^{-1} \quad \mathbf{U}^{-1} \mathbf{A} \mathbf{V} = \mathbf{\Sigma} = \begin{bmatrix} \sigma_1 & & & \\ & \sigma_2 & & \\ & & \ddots & \\ & & & \sigma_N \end{bmatrix}$$

Singular Value Decomposition

- SVD is ALWAYS possible on a real matrix A
- U , Σ , V unique
- U and V are column orthonormal
 - $U^T U = I$, $V^T V = I$
 - Columns are orthogonal unit vectors
- Σ : Diagonal
 - Entries are nonzero and sorted in descending order

Singular Value Decomposition

```
U, S, V_transpose = np.linalg.svd(M)
```

```
print(U)
```

```
[[-0.95123459  0.23048583 -0.20500982]
 [-0.28736244 -0.90373717  0.31730421]
 [-0.11214087  0.36074286  0.92589903]]
```

```
print(S)
```

```
[ 3.72021075  2.87893436  0.93368567]
```

```
print(V_transpose)
```

```
[[-0.9215684  -0.03014369 -0.38704398]
 [-0.38764928  0.1253043   0.91325071]
 [ 0.02096953  0.99166032 -0.12716166]]
```



$$M = \begin{bmatrix} 3 & 0 & 2 \\ 2 & 0 & -2 \\ 0 & 1 & 1 \end{bmatrix}, \quad v = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

Recall SVD is the factorization of a matrix into the product of 3 matrices, and is formulated like so:

$$M = U\Sigma V^T$$


Caution: The notation of SVD in NumPy is slightly different. Here V is actually V^T in the common notation.

Dimensionality Reduction



A fun fact

SVD provides the best rank-k approximation of A (in terms of reconstruction error)!



In other words: B is a solution to $\min_B \|A - B\|_F$

More Information

Python Documentation: <https://docs.python.org/3/>

NumPy Documentation: <https://numpy.org/doc/2.2/reference/index.html>

CS 229 Linear Algebra Tutorial: : <https://cs229.stanford.edu/section/cs229-linalg.pdf>

CS231N Python Tutorial: <http://cs231n.github.io/python-numpy-tutorial/>

Office hours! Ed! The Internet!

PSET0 Overview

Announcements

- Lecture on Mon and Wed virtual – TAs in person!
- PSET 1 Released!
- Project partner searching thread on Ed

Thanks!

Questions?